

# Lecture 10 – Using Thread Pools

- **Housekeeping:**
  - **exam return Monday**
  - **Programming assignment on Monday**
    - **Purpose: practice use of the mechanisms we've discussed**
    - **The next assignment will be more oriented toward system design using concurrency**
    - **Implement the synchronizers of Section 5.5**

# Quickly, task cancellation

- Chapter 7, not covering in detail
- Only a few observations on cancellation
  - Why? Nobody cares about the work a thread is doing any more – taking too long, some other thread got the answer, another thread got an error that is fatal for the app as a whole
  - How? Cooperatively! Much saner than C signal (interrupt) handling
    - Thread handles cancellation only when in known state
    - C program may have to handle a signal in any state
  - Beware: Intrinsic locks don't respond to interruption
  - Daemon thread: in Java a thread whose existence

# Chapter 8 – Applying Thread Pools

- **A lot of Java-specific details in this chapter**
- **Lecture will try to illuminate the underlying issues a bit more**
  - **Client dependence on execution policy**
  - **Thread-pool sizing**
  - **Queuing policy**
  - **Saturation policy**
  - **Thread-creation**

# Client dependence on execution policy

- **Tasks that use objects that rely on thread-confinement for safety**
  - Recall the single-threaded executor
- **Inter-dependent tasks – ones that rely on state changes or results produced by other tasks**
  - Avoiding thread-starvation deadlock
  - Timing-sensitive tasks – a less dramatic form of starvation

# Thread-starvation deadlock

- **Task A, currently executing in a thread, requires results from task B**
- **Task B cannot be executed because there are no available threads in the thread pool**
- **Obvious possibility in a single-thread executor but can happen with any executor that has a maximum number of threads**
  - **Consider: is the situation any better using Thread directly?**

# You can have more than one thread pool

- **Suggestion: in any given thread-pool have *homogeneous* tasks – similar lifetimes, similar work.**
  - **It is hard to choose thread-pool policies if arbitrary tasks may be thrown at the pool**
    - Long-running tasks keep short-running tasks from executing
    - Priority between vastly dissimilar tasks may be hard to pin down
  - **Neither thread pools nor threads are so expensive that you have to worry very much about whether you have 100 or 200**
    - You may have to worry about the difference between 100 and 1000
- **Remember that once a thread pool assigns a task to a**

# Thread-pool sizing

- **Three thread states: waiting, ready-to-run, running**
  - **Waiting threads main effect is that they take up memory**
  - **Ready-to-run threads compete for the processor; lots of ready threads may make each scheduling decision more costly**
  - **Running threads – limited to the number of processors**
  - **Switching between ready and run states is**

## Pool sizing (2)

- The following are all system-wide (not per thread-pool) considerations
- Too many waiting threads is bad
  - How many is too many?
- Too many ready threads is bad
  - How many is too many?
  - Better question: how many is enough?
- Answer: too few running threads is bad, so ideally # of ready threads is 0 and number of running threads = number of CPUs



## Pool sizing (3)

- **For a resource-limited pool consider**
  - **Characteristics of tasks it will service**
    - Task execution time
    - Task wait time
  - **and how many resources should be devoted to those tasks**
    - What else does the system have to do concurrently?
  - **$N_{\text{threads}} = N_{\text{cpu}} * U * (1 + W/C)$** 
    - $U$  == desired CPU utilization for the whole pool
    - $W$  == task wait time.  $C$  == task compute time

## Pool sizing (4)

- $\text{Nthreads} * \text{per-task-disk-utilization} = \text{desired pool-disk-utilization}$
- $\text{Nthreads} * \text{per-task-network-utilization} = \text{desired pool-disk-utilization}$
- Etc.
- $\text{Pool size} = \text{minimum of all of the above} - \text{it doesn't help to make it any bigger}$

## Pool sizing (5)

- **Need to be concerned also with the size of the queue “in front of” the thread pool**
  - **Implicit assumption: a task can be represented by much less data than a thread**
    - **Number of waiting tasks >> number of threads**
    - **On the other hand...**
  - **What are queuing systems good for?**
    - **Smoothing out variations in arrival rate and service time**
    - **If work continuously arrives at a faster rate than it can be performed then queue length and wait times**

# Queuing Policy

- **Fundamental question: bounded or unbounded queue**
  - **Unbounded queue allows memory exhaustion and unbounded wait times**
  - **Bounded queue raises question of what to do when it is full – the *saturation policy***
    - **Abort – throw an exception**
    - **Discard – silently ignore submitted task**
    - **Discard oldest – or other policy that throws away some already queued task**
    - **Caller runs – run the task in the submitting thread**

# Work rejection by the system

- In the exam problem work is internally generated by the application. Conceptually easy to slow down work submission
- What if work is being delivered by an external source, e.g. over the network – need to think about how back pressure can be exerted on the work sources
  - Classic denial of service attack – no effective way to exert such backpressure
  - Other kinds of DoS may work by breaking

# Discussion

- **Given the dangers of unbounded thread pools and unbounded queues why is an unbounded queue the default for fixed size thread pools and why do `cachedThread` pools support an unlimited number of threads?**

# Thread Creation

- **Thread pools need to create threads from time to time**
- **A default thread Factory is supplied but you can customize the thread pool with your own factory (what is a factory?)**
  - **Statistics gathering**
  - **Customized unhandled exception processing**
  - **Customized security policies**
  - **...**
- **All this can be done without using a custom factory without using a thread pool at all but a thread pool provides a nice point of focus**